

Assignment 4 & 5

Code Optimization in C:

The Gravitational N-Body Problem Using the The Barnes-Hut Algorithm

Finn Joel Bjervig*, August Forsman†, Maja Linderholm‡ and Erik
Turesson§

*1TD062 - High Performance Programming,
Department of Information Technology,
Uppsala university, Sweden*

August 31, 2022

*Electronic address: joelfbjervig@gmail.com

†Electronic address: aufo8456@student.uu.se

‡Electronic address: maja.linderholm@gmail.com

§Electronic address: ertu2293@student.uu.se

The Problem

Constructing simulations is at first glance a task of implementing equations and numerical methods to approximate solutions, up until the programs performance heavily rely on the scalability of its algorithms. Different methods scale differently, and even arithmetic operations in the computer can be chosen with precaution to make the program run faster. This assignment aims to simulate a closed, dynamical system of particles of varying mass and population (from a couple of particles up to thousands), subject to the exerted gravitational forces by each other. The gravitational force exerted on particle i by a particle j is

$$\mathbf{f}_{ij} = -G \frac{m_i m_j}{r_{ij}^2} \hat{\mathbf{r}}_{ij}$$

where m_i , m_j are the masses of the particles, G is a gravitational constant, and \mathbf{r}_{ij} is the vector from i to j . In this project, we consider a system of N particles with $G = 100/N$ and a slight modification in that we assume a minimal possible distance $\epsilon_0 = 10^{-3}$ between particles. Thus, each particle i is affected by the total force

$$\mathbf{F}_i = -G m_i \sum_{j=0, j \neq i}^{N-1} \frac{m_j}{(r_{ij} + \epsilon_0)^3} \hat{\mathbf{r}}_{ij}$$

The only analytical solution of such a system is limited to only two particles, well known as the two body problem (except for very specific cases of the equally famous three body problem). As soon as one introduces another body of significant mass, the system becomes chaotic. Computers have the advantage of numerical approximations but at a scale of thousands of particles the necessity of writing efficient code cannot be overstated. As the reader will see, there are many ways to bring down the running time of such a program, some of which can be used in general when coding, and other tricks which is acknowledged in this specific task.

The intuitive way of implementing the description above is to simply iterate through all particles, and calculate all forces asserted by all other particles. This *straightforward algorithm* (for lack of a better name), does yield a correct and accurate result, but is quite slow. Simulations and computational science boils down to a balance between accuracy and efficiency, and by acknowledging that compromise, one may look to algorithms that are far superior in speed but may not be as accurate as its competitor. For N-body simulations there exists an approximate solution called *The Barnes-Hut Algorithm*. The idea is to, in a quadrant, approximate a cluster of particles with a shared center of mass, and the total mass of all particles. Thereby it can compare one particle with this collective virtual particle and doesn't need to calculate all forces between all particles. Depending on the size of the cluster and the distance between the cluster and the particle, the approximation is more or less valid. To determine whether to proceed with the calculation or split up the cluster further, the *Theta Criterion* is evaluated. A constant θ_{max} is set by the user which is defined as

$$\theta = \frac{\text{side length of cluster square frame}}{\text{distance from particle to center of cluster frame}}$$

If $\theta < \theta_{max}$ there's sufficient precision and the current cluster can be used, otherwise the cluster is recursively split into smaller quadrant clusters until the condition is fulfilled.

The Solution

To calculate the evolution of particles N with a "straightforward"-algorithm and Barnes Hut algorithm in a gravitational simulation the code, which is listed in the Appendix, is structured in the following way.

Straightforward Algorithm

Data structures

Each particle is defined as a **struct** containing the mass, brightness, and vectors for position *pos*, velocity *vel*, and acceleration *acc*. Using only arrays to store the values of each particle did increase the performance but not significantly. Therefore, the less cluttered **struct** implementation was chosen.

Functions

main: Driver function of the simulation. Handles input arguments and graphics.
read_doubles_from_file: Reads initial data from the input file for each particle and stores it in a struct. An input file contains the initial position, velocity, mass and brightness for each particle.

time_sim_vec: Simulates the movement of the particles over time. The function will loop over each particle and calculates the force it is subjected to by other particles.

write_to_file: Writes results of final time step to output file.

Barnes Hut Algorithm

Data structures

The struct is implemented as in the straightforward algorithm. The N -particle structs are inserted into a quad-tree, which is a tree structure where each node has four children. It is essentially an adaption for a binary tree to store 2-dimensional data. In the root node, the domain of computation is subdivided into four fields (using notations NW, NE, SW, SE for each intercardinal direction) that can be accessed through a corresponding pointer. A node can be visualized as a rectangular domain and has the ability to split into four children according to the Barnes-Hut algorithm. When a layer is added, the node struct parameters **length** and **pos** are initialized from the values of its parent. Once a

particle is inserted into a node, the `total_mass`, and `center_mass` gets updated correspondingly and enables the program to compute forces from multiple node layers using a cluster of particles in the 2D-plane.

Functions

`main`: Driver function of the simulation. Handles input arguments and graphics.
`read_file`: Reads initial data from the input file for each particle and stores it in a struct. An input file contains the initial position, velocity, mass and brightness for each particle.

`write_to_file`: Writes results of final time step to output file.

`barnes_hut`: Simulates the movement of the particles over time. The function will loop over each particle and calculates the force it is subjected to by other particles. Calls functions `init`, `insert`, `force_traverse`, `del_tree` and `tree_draw`.

`init`: Initializes the root node of the quadtree.

`insert`: Inserts particles in the quadtree recursively.

`force_traverse`: Calculates the acceleration of each particle.

`del_tree`: Deletes the quadtree.

`tree_draw`: Draws the tree.

Optimizations

Variable Keywords

By using the `register` keyword the compiler tries to keep the variable in one of the CPU-registers instead of the memory. This minimizes memory accesses and optimizes the code on an instruction level. It was used on the most commonly used variables, such as the loop variables. Similarly, `const` was applied inside the time stepping function to applicable variables such as `N`, `delta_t`, and `G`. Together, these keywords reduced the execution time by almost 30% when no compiler optimization flags were being used. However, this did not seem to have a noticeable effect for more aggressive flags.

Vectorization

The vector datatype `vector_t` is defined as two allocated memory slots of size 2×8 bytes (corresponding to the size of two `double`) such that they're positioned next to each other. This memory alignment enables for arithmetic operations between such vectors, or vectors and scalars (piece-wise operations), to be computed during the same clock cycle. This method of *vectorizing* the code improves the running time significantly.

Calculating Distances

By storing θ_{max}^2 instead of just θ_{max} in our Barnes-Hut implementation, we can modify the condition to $\theta_{max}^2 \geq side^2/distance^2$. This means that we avoid

having to calculate the root of the distance which is a very costly operation. For larger N 's, this alone reduced the execution time of our Barnes-Hut implementations by almost 25%, as can be seen in Figure 1.

Initialization of Child Nodes

When particles are inserted in the quadtree nodes are often split. When a split occurs, the former leaf node becomes a parent node and two different ways of initializing the new child nodes was implemented. Either all child nodes (NW, NE, SW, SE) was initialized even though only e.g. one of them was used to store a particle, or only the child node used to store the particle was initialized. Figure 1 shows the execution time for the two implementations and concludes that initializing only relevant child nodes yields faster execution.

Optimal θ_{max} for Barnes-Hut

In order to find a suitable value of θ_{max} we wanted to maximize it for higher performance while maintaining a target accuracy of 10^{-3} . The accuracy is measured as the maximal positional difference for any particle between its solutions in the exact and the Barnes-Hut algorithm after a certain amount of time steps. In table 1 we list measurements of this accuracy for some values of θ_{max} .

θ_{max}	Accuracy
1	0.0163
0.75	0.0122
0.5	0.0056
0.4	0.0047
0.3	0.0023
0.25	0.0009
0.0	0.0000

Table 1: The maximal positional difference between the exact and the Barnes-Hut solution after 200 time steps depending on the value of θ_{max} . These results were produced from a system of 2000 particles with $\Delta t = 10^{-3} s$

As one can see, $\theta_{max} = 0.25$ just about reaches the target accuracy.

Performance and Discussion

The simulations ran on a *11th Gen Intel(R) Core(TM) i7-1165G7 @ 2.80GHz* CPU, compiled with GCC 10.2.0 on Ubuntu 20.04 through wsl (Windows Subsystem for Linux). We used `-Ofast`, `-march=native`, `-ffinite-math-only`, and `-fno-signed-zeros` as compiler flags to optimize for performance.

In figure 1 below, the execution times for our implementation times can be seen. For the Barnes-Hut algorithms, $\theta_{max} = 0.5$ was used. As can be seen,

the Barnes-Hut algorithms scale as $\mathcal{O}(n \log n)$ when increasing the number of particles and run much faster compared to the $\mathcal{O}(n^2)$ algorithm.

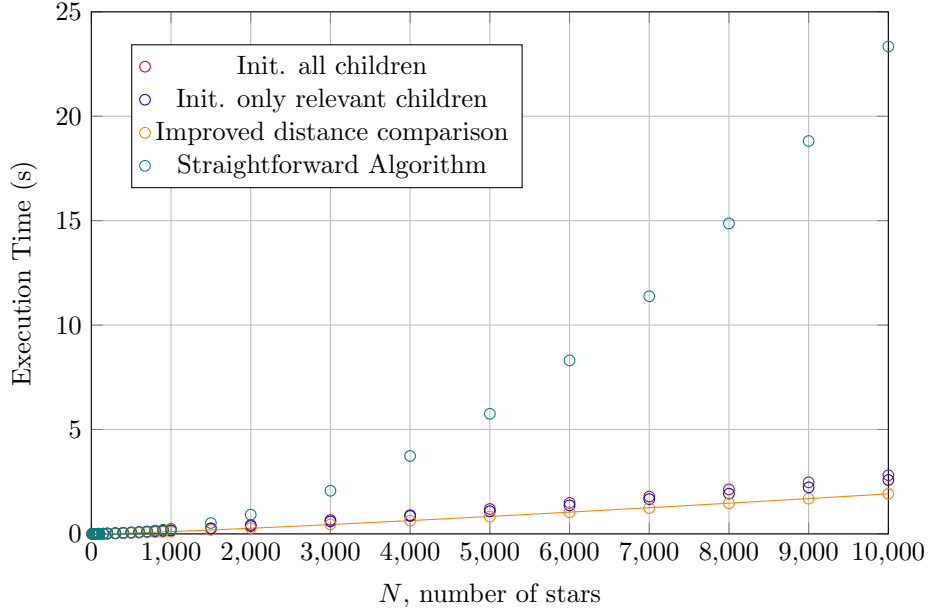


Figure 1: Serial execution times for 100 time steps using our optimizations for the $\mathcal{O}(n \log n)$ Barnes-Hut algorithm compared to our implementation of an $\mathcal{O}(n^2)$ straightforward algorithm. $\theta_{max} = 0.5$ was used for the Barnes-Hut executions. The solid line is a fitted $N \log N$ approximation of the best Barnes-Hut implementation.

Parallelization using Pthreads

Splitting up the work to several threads enables for tasks to be computed in parallel. The threads are then distributed between the CPU's cores. The number of cores and the number of threads per core depends on the CPU model, but many more recent personal computers will usually have 4-16 cores and 2 threads per core. As previously mentioned, our test were ran on a *11th Gen Intel(R) Core(TM) i7-1165G7 @ 2.80GHz* CPU, which has 4 cores with up to 2 threads per core.

If implemented perfectly, the speedup of an embarrassingly parallel problem should be directly proportional to the number of threads used. Two threads would then correspond to a double speedup, and three threads to a three-fold speedup, etc.

Parallelization of the Straightforward Algorithm

When examining our straightforward implementation, we found that we spent around 98% of the time in calculating the forces for problems of large N s. Calculating the forces for particles can be done individually for each particle and should thus be suitable to parallelization. The amount of work when iterating through the particle array is always the same: $N - 1$ computations for each of the N particles (the scaling of this method is $\mathcal{O}(N(N - 1))$). The optimal splitting of tasks is therefore an even split such that each thread manages N/P particles, where P is the number of threads used.

Figure 2 displays the *straightforward* algorithm's performance depending on the number of threads. Interestingly, beyond around four threads, there is not much of a difference performance-wise, no matter how many threads is used. This may be explained by the fact that the threads don't work independently. The particles handled by one thread still needs the forces exerted by particles in the other threads, which causes delay. Additionally, it might be due to the fact that while the CPU used can handle 2 threads per core, it still just has 4 cores. Going above this limit might be the reason for the performance decrease in scaling.

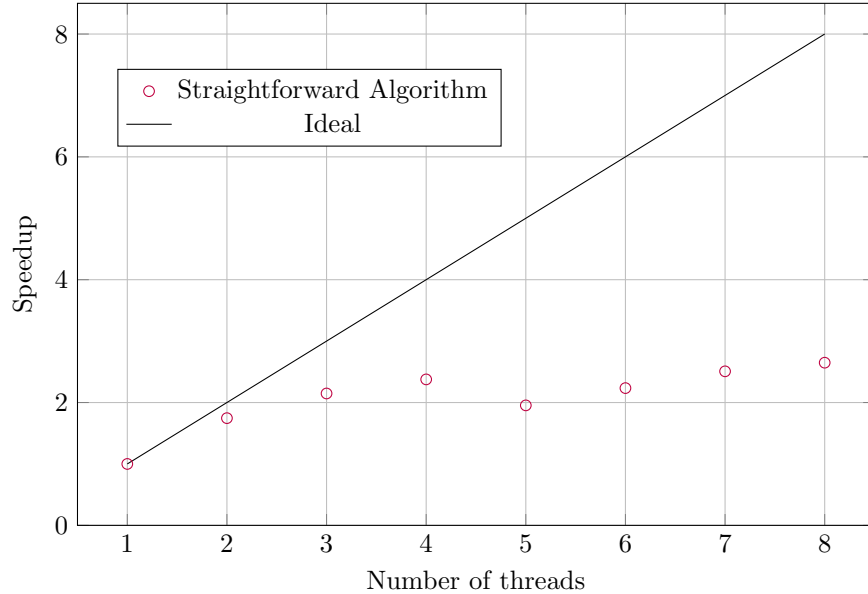


Figure 2: The relative speedup of our Straightforward implementation versus the number of threads used

Originally, we tried another approach where we used Newton's third law such that when particle i has its subjected forces computed by particle j , the latter will also be assigned the equal but opposite force. The scaling is therefore half

of the previous one, $\mathcal{O}(\frac{1}{2}N(N-1))$. While the serial implementation of this approach had about half the execution time, it showcased very bad scaling due to uneven workload between threads, and a lot of overhead. Thus we chose to not go further with implementing that solution in this assignment. However, below one can find a discussion of how workload could theoretically be more evenly divided. This splitting was not implemented but can be investigated for further work.

If we use Newton's third law, even if we would have the same number of particles in an evenly split the particle array, the number of computations in those intervals will not be equally many. In the beginning of the array there are $N-1$ forces to compute, then $N-2$ until the last particle where there are no forces left to compute. Therefore, the amount of work decreases linearly from N to 0 as: *number of computations* = $N - \text{particle array index}$. In figure 3b the four intervals are set such that the total work for each thread (the area under the curve, A_i), is the same for all threads. See Figure 3a, for an example of splitting between four threads.

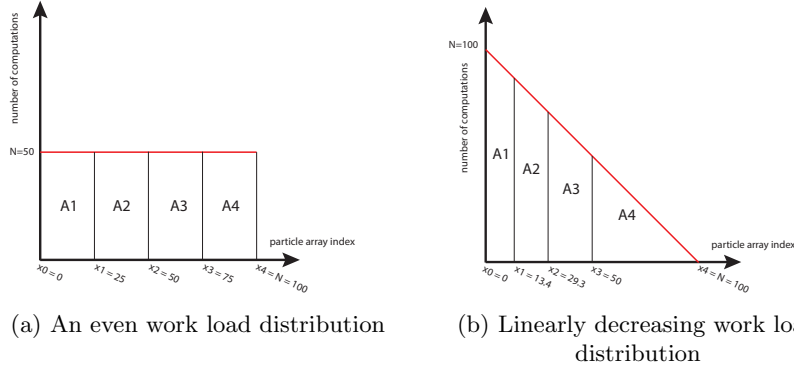


Figure 3: On the premise of equal work among the threads, $A_1 = A_2 = A_3 = A_4$ the number of particles handled by each thread is different for the two scenarios

The bounds are derived from the fact that $A_1 = A_2 = A_3 = A_4$ and that $P \cdot A = \frac{1}{2}N^2 \quad \forall A_i, \quad i \in [1, 4]$. This yields the following equation

$$x_{i+1} = -\sqrt{(x_i - N)^2 - \frac{N^2}{P}} + N \quad (1)$$

where N is the number of particles in the system, P are the number of threads used and the initial point is $x_0 = 0$.

Parallelization of the Barnes-Hut Algorithm

For the parallelization of our Barnes-Hut Algorithm we started at looking what parts of the code took the longest time for large values of N . For our serial implementation of Barnes-Hut, we found that for a simulation of 100 time steps of 20000 particles, using $\theta_{max} = 0.25$, we spent about 96.5% of the total execution time calculating and updating the forces for each particle, 3.4% for building and deleting the quadtree at each time step, and the rest for file I/O. Thus, we chose to spend our effort in trying to parallelize only the force calculations.

Our multithreaded implementation utilizes one master thread that handles the construction of the quadtree at each time step. Once the tree is built, the master thread reaches a barrier. Worker threads wait at this barrier, and then starts updating the particles. Noting that the force calculation should be independent for each particle, we evenly divided the particles between the worker threads. Once all threads have updated all its particles, the main thread is unlocked, and deletes the tree. This process is then repeated until we have reached the desired number of time steps.

As seen in figure 4, the speedup is close to the ideal (1.85 for 2 threads) in the beginning, but stagnates with increasing number of threads. Important to note is however that the speedup varies significantly between different θ_{max} , which implies the importance of the parameter beyond the scaling of the number of particles.

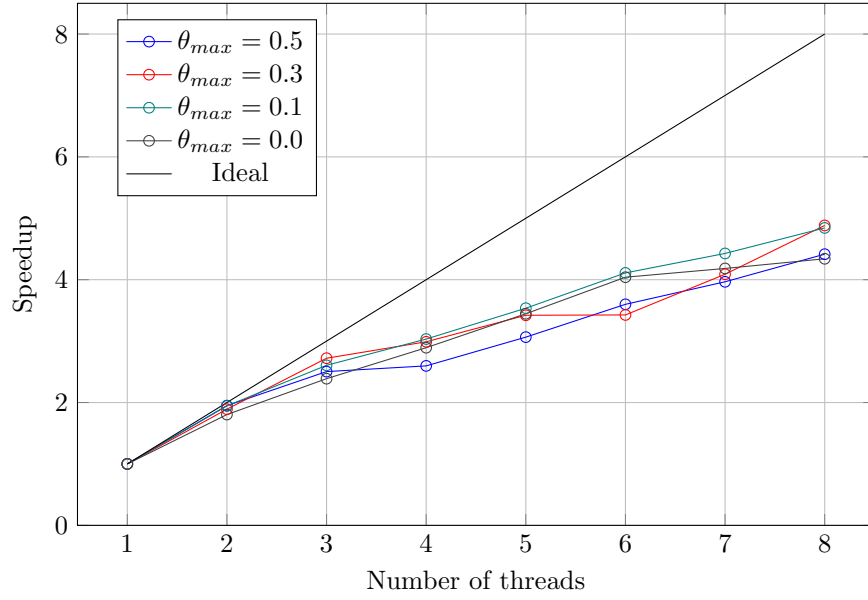


Figure 4: The relative speedup of our Barnes-Hut implementation versus the number of threads used for some different values of θ_{max}

When using tree structures as in the Barnes-Hut algorithm it is difficult to know how to split up the load between the threads, since some particles will require more recursive descent than others in order to fulfill the *Theta criterion*. We suspected the stagnant scaling to be due to uneven load-balancing between the threads. Therefore, we tried another queue-based approach where threads took a particle index from a global counter, starting at $N - 1$, before decrementing it until it reached 0. While this approach seemed to display a somewhat better scaling, the baseline performance was a lot worse, and thus we chose to not explore that possibility further.

It is evident in figure 4 that the program benefited from parallelization. The decreased slope of the data points can be derived from Amdahl's law

$$\frac{1}{1 - p} = 4.88 \quad (2)$$

where 4.88 is the highest measured speedup for 8 threads. Solving $p \approx 0.80$ means that the implemented techniques was able to parallelize 80% of the program.

Appendix

Serial Straightforward galsim.c

```
1 #include <math.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include "graphics.h"
5
6 typedef double vector_t __attribute__((vector_size(2 * sizeof(
    double))));
7
8 #define EPSILON_0 1e-3
9
10 typedef struct particle
11 {
12     /*
13      * Particle properties:
14      * vel - velocity
15      * pos - coordinates
16      */
17     vector_t pos;
18     double mass;
19     vector_t vel;
20     double brightness;
21     vector_t acc;
22 } p_t;
23
24 void read_doubles_from_file(int N, p_t *p, const char *fileName);
25 void print_struct(p_t p);
26 //void time_sim(int nsteps, const double delta_t, p_t *p, int N,
    const int graphics);
27 void write_to_file(p_t *p, int N, char *fileName);
28 //void time_sim2(int nsteps, const double delta_t, p_t *input_array
    , int N, const int graphics);
29 void time_sim_vec(const int nsteps, const double delta_t, p_t *p,
    const int N,
30                  const int graphics);
31
32 int main(int argc, char *argv[])
33 {
34     // Check correct number of input arguments
35     if (argc != 6)
36     {
37         printf("Wrong number of inputs \n");
38         exit(EXIT_FAILURE);
39     }
40
41     // Define input variables
42     const int N = atoi(argv[1]);
43     const int nsteps = atoi(argv[3]);
44     const double delta_t = atof(argv[4]);
45     const int graphics = atoi(argv[5]);
46
47     p_t p[N];
48
49     if (graphics)
```

```

50 {
51     InitializeGraphics(argv[0], 800, 800);
52     SetCAxes(0, 1);
53     read_doubles_from_file(N, p, argv[2]);
54
55     // print structs if wanted
56     //for (int i=0; i<N; i++) print_struct(p[i]);
57     time_sim_vec(nsteps, delta_t, p, N, graphics);
58     FlushDisplay();
59     CloseDisplay();
60     write_to_file(p, N, "result.gal");
61 }
62 else
63 {
64     read_doubles_from_file(N, p, argv[2]);
65
66     // print structs if wanted
67     //for (int i=0; i<N; i++) print_struct(p[i]);
68     time_sim_vec(nsteps, delta_t, p, N, graphics);
69     write_to_file(p, N, "result.gal");
70 }
71 return 0;
72 }
73
74 void write_to_file(p_t *p, int N, char *fileName)
75 {
76     FILE *file = fopen(fileName, "wb");
77
78     for (int i=0; i<N; i++){
79         fwrite(&p[i].pos,      sizeof(p[i].pos),      1, file);
80         fwrite(&p[i].mass,     sizeof(p[i].mass),      1, file);
81         fwrite(&p[i].vel,      sizeof(p[i].vel),      1, file);
82         fwrite(&p[i].brightness, sizeof(p[i].brightness), 1, file);
83     }
84     fclose(file);
85 }
86
87 void print_struct(p_t p)
88 {
89     printf(
90         "x: %lf\t y: %lf\t mass: %lf\t vel_x: %lf\t vel_y: %lf\t\n",
91         p.pos[0], p.pos[1], p.mass, p.vel[0], p.vel[1], p.
92         brightness);
93 }
94
95
96 ///////////////////////////////////////////////////
97 // READ DOUBLES //
98 ///////////////////////////////////////////////////
99 // NOTE: might not work because vector_t might not operate the same
100 // way as a regular variable.
101
102 void read_doubles_from_file(int N, p_t *p, const char *fileName)
103 {
104     /* Open input file and determine its size. */

```

```

104 FILE *file = fopen(fileName, "rb");
105 if (!file)
106 {
107     printf("read_doubles_from_file error: failed to open input
108     file '%s'.\n",
109         fileName);
110     exit(EXIT_FAILURE);
111 }
112 // zerovector for vectorization convenience
113 vector_t zerovec;
114 zerovec[0] = 0;
115 zerovec[1] = 0;
116
117 p_t temp;
118 double buffer[6]; // file reading buffer
119 for(int i = 0; i < N; i++) {
120     // Read doubles into particle struct
121     fread(&buffer, sizeof(double), 6, file);
122     temp.pos[0] = buffer[0];
123     temp.pos[1] = buffer[1];
124     temp.mass = buffer[2];
125     temp.vel[0] = buffer[3];
126     temp.vel[1] = buffer[4];
127     temp.brightness = buffer[5];
128     temp.acc = zerovec;
129     p[i] = temp;
130 }
131 // Close file and check if closed successfully
132 if (fclose(file) != 0)
133 {
134     printf("read_doubles_from_file error: error closing input
135     file.\n");
136     exit(EXIT_FAILURE);
137 }
138
139 // VECTORIZED
140 void time_sim_vec(const int nsteps, const double delta_t, p_t *p,
141     const int N,
142     const int graphics)
143 {
144     vector_t diff, diffsqrd, F;
145
146     // zerovector for vectorization convenience
147     vector_t zerovec;
148     zerovec[0] = 0;
149     zerovec[1] = 0;
150
151     double distance;
152     const double G = 100.0 / N;
153     float circleradius[N], circlecolor = 0;
154
155     if (graphics)
156     {
157         // make radius dependent on the mass of the particle
158         for (register int i = 0; i < N; ++i)
159         {

```

```

158     circleradius[i] = p[i].mass * 1.5e-3;
159 }
160 }
161 // loop over all timesteps
162 for (int t = 0; t < nsteps; t++)
163 {
164
165     // reset all p's accelerations to zero for each time step
166     for (register int i = 0; i < N; i++)
167     {
168         p[i].acc = zerovec;
169     }
170
171     // iterate through all particles
172     for (register int i = N-1; i >= 0; i--)
173     {
174         for (register int j = i - 1; j >= 0; j--) // only up til
the particle in the above for loop
175         {
176             // temp variables
177             diff = p[i].pos - p[j].pos; // particles relative position
in x,y coordinates
178             diffsqrd = diff*diff; // - squared
179             distance = sqrt(diffsqrd[0]+diffsqrd[1])+ EPSILON_0; //
pythagoras thrm + small number for
180 //
smoothing when distance is small
181
182             // NOTE: maybe faster to just have everything on one long
calculations,
183             // because now we store them in temporary variables so its
easier to see whats happening
184             // something we could do in the end?
185             F = G*diff/(distance*distance*distance); // 1/r^2 in the
direction of r
186             p[i].acc +=p[j].mass*F;
187             p[j].acc -=p[i].mass*F;
188         }
189     }
190     // when all the particles accelerations have been calculated,
191     // we can now, for the same time t_i update all particles
velocities
192     // and positions, so that no particle is in a different
timestep
193     for (register int i = 0; i < N; i++)
194     {
195         p[i].vel -=delta_t* p[i].acc;
196         p[i].pos += delta_t*p[i].vel;
197     }
198
199     if (graphics)
200     {
201         ClearScreen();
202         // picasso away!
203         for (int i = 0; i < N; i++) DrawCircle(p[i].pos[0], p[i].pos
[1], 1, 1, circleradius[i], circlecolor);
204         Refresh();

```

```

205     usleep(30);
206 }
207 }
208 }

```

Serial Barnes-Hut galsim.c

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <math.h>
4  #include "graphics.h"
5
6  #define EPSILON_0 1e-3    //Plummer radius
7  #define W 1              // Window width
8  #define L 1              // Window length
9
10
11 typedef double vector_t __attribute__((vector_size(2 * sizeof(
12     double)))));
13
14 /*
15  * Particle:
16  * pos — coordinates
17  * mass — mass of particle
18  * vel — velocity
19  * brightness — brightness of particle
20  * acc — acceleration
21 */
22 typedef struct particle
23 {
24     vector_t pos;
25     double mass;
26     vector_t vel;
27     double brightness;
28     vector_t acc;
29 } p_t;
30
31 /*
32  * Quad tree node:
33  * particle — particle struct pointer
34  * length — length of box/quadnode
35  * total_mass — mass of all particles inside a box
36  * center_mass — position of center of mass
37  * pos — coordinate of center of box
38  * NE,NW,SW,SE — quadrants of domain. NE = northeast, etc.. (
39     these are the nodes children)
40 */
41 typedef struct quad_tree_node
42 {
43     p_t *particle;
44     double length;
45     double total_mass;
46     vector_t center_mass;
47     vector_t pos;
48     struct quad_tree_node *NE;
49     struct quad_tree_node *NW;

```

```

48     struct quad_tree_node *SW;
49     struct quad_tree_node *SE;
50
51 } q_node_t;
52
53 /*
54     Initialize global helpfull vectors
55 */
56 const vector_t zerovec = {0.0, 0.0};
57 const vector_t se_vec = {1.0, 1.0};
58 const vector_t sw_vec = {-1.0, 1.0};
59 const vector_t ne_vec = {1.0, -1.0};
60 const vector_t nw_vec = {-1.0, -1.0};
61
62
63
64 /*
65 Macro function that creates a child to the parent, with a vector
66 in the corresponding direction.
67     Example: Create a child to the north-west to a q_node_t ** node
68     create_child((*node)->NW, (*node), nw_vec);
69 */
70 #define create_child(child, parent, dir) \
71 ({ \
72     child = (q_node_t *) malloc(sizeof(q_node_t)); \
73     child->pos = parent->pos + 0.25 * (parent->length) * dir; \
74     child->length = parent->length * 0.5; \
75     child->total_mass = 0.0; \
76     child->particle = NULL; \
77     child->NW = NULL; \
78     child->NE = NULL; \
79     child->SW = NULL; \
80     child->SE = NULL; \
81 })
82
83
84 /*
85 Macro function that locates the correct child quadrant to insert a
86 particle
87 into. If val==0: check if child already exists. If not, create it
88 before inserting.
89 If val==1: insert knowing that child does not exist.
90 */
91 #define split_insert(val, particle) \
92 ({ \
93     if ((particle->pos[0] >= (*node)->pos[0]) && (particle->pos[1] \
94     < (*node)->pos[1])) { \
95         if (val || (*node)->NE == NULL) { \
96             create_child((*node)->NE, (*node), ne_vec); \
97         } \
98         insert(&((*node)->NE), particle); \
99     } \
100     else if ((particle->pos[0] < (*node)->pos[0]) && (particle->pos \
101     [1] < (*node)->pos[1])){ \
102         if (val || (*node)->NW == NULL) { \
103             create_child((*node)->NW, (*node), nw_vec); \
104         } \
105     } \

```



```

101         insert(&((*node)->NW), particle); \
102     } \
103     else if ((particle->pos[0] < (*node)->pos[0]) && (particle->pos
104     [1] >= (*node)->pos[1])) { \
105         if (val||(*node)->SW == NULL) { \
106             create_child((*node)->SW, (*node), sw_vec); \
107         } \
108         insert(&((*node)->SW), particle); \
109     } \
110     else { \
111         if (val||(*node)->SE == NULL) { \
112             create_child((*node)->SE, (*node), se_vec); \
113         } \
114         insert(&((*node)->SE), particle); \
115     } \
116 })
117
118
119 /*
120     Declare functions
121 */
122
123 //IO functions
124 void read_file(int N, p_t *p, const char *fileName);
125 void write_to_file(p_t *p, int N, char *fileName);
126 void print_struct(p_t p); //mainly for debugging
127
128 q_node_t *init(); //create tree root
129 void insert(q_node_t **node, p_t *particle); //recursive insertion
130     of particle into tree
131 void tree_del(q_node_t **node); //free memory recursively
132 void barnes_hut(const int nsteps, const double delta_t, p_t *p,
133     const int N,
134     const double theta_max, const int graphics); //the
135     time stepping algorithm
136 void force_traverse(q_node_t **node, const double theta_max, p_t *
137     particle); //calculation of the forces
138
139 void tree_draw(q_node_t **node); //recursive drawing of the
140     quadnodes
141
142
143 /*
144     Driver function
145 */
146 int main(int argc, char *argv[])
147 {
148     // Check correct number of input arguments
149     if (argc != 7)
150     {
151         printf("Wrong number of inputs \n");
152         exit(EXIT_FAILURE);
153     }
154
155     // Define input variables

```

```

152     const int N = atoi(argv[1]);
153     const int nsteps = atoi(argv[3]);
154     const double delta_t = atof(argv[4]);
155     const double theta_max = atof(argv[5]) * atof(argv[5]); //
156     square for cheaper comparison in the force calculation
157     const int graphics = atoi(argv[6]);
158
159     //Read from input file and save in an array of particles
160     p_t p[N];
161     read_file(N, p, argv[2]);
162
163     // initialize window if graphics is enabled
164     if (graphics) InitializeGraphics(argv[0], 800, 800);
165
166     //Simulate particle system with Barnes Hut algorithm and save
167     result
168     barnes_hut(nsteps, delta_t, p, N, theta_max, graphics);
169     write_to_file(p, N, "result.gal");
170     return 0;
171 }
172
173 /*
174  Function that reads doubles from input file and stores them in
175  an array of structs
176 */
177 void read_file(int N, p_t *p, const char *fileName)
178 {
179     /* Open input file and determine its size. */
180     FILE *file = fopen(fileName, "rb");
181     if (!file)
182     {
183         printf("read_doubles_from_file error: failed to open input
184         file '%s'.\n",
185             fileName);
186         exit(EXIT_FAILURE);
187     }
188
189     p_t temp;
190     double buffer[6]; // file reading buffer
191     for (int i = 0; i < N; i++)
192     {
193         // Read doubles into particle struct
194         if (0==fread(&buffer, sizeof(double), 6, file)) exit(
195         EXIT_FAILURE);
196         temp.pos[0] = buffer[0];
197         temp.pos[1] = buffer[1];
198         temp.mass = buffer[2];
199         temp.vel[0] = buffer[3];
200         temp.vel[1] = buffer[4];
201         temp.brightness = buffer[5];
202         temp.acc = zerovec;
203         p[i] = temp;
204     }
205     // Close file and check if closed successfully
206     if (fclose(file) != 0)

```

```

204 {
205     printf("read_doubles_from_file error: error closing input
206           file.\n");
207     exit(EXIT_FAILURE);
208 }
209
210 /*
211  Function that writes output to file
212 */
213 void write_to_file(p_t *p, int N, char *fileName)
214 {
215     FILE *file = fopen(fileName, "wb");
216
217     for (int i = 0; i < N; i++)
218     {
219         fwrite(&p[i].pos, sizeof(p[i].pos), 1, file);
220         fwrite(&p[i].mass, sizeof(p[i].mass), 1, file);
221         fwrite(&p[i].vel, sizeof(p[i].vel), 1, file);
222         fwrite(&p[i].brightness, sizeof(p[i].brightness), 1, file);
223     }
224     fclose(file);
225 }
226
227 /*
228  Function that prints values of a struct. Used for debugging
229 */
230 void print_struct(p_t p)
231 {
232     printf(
233         "x: %lf\t y: %lf\t mass: %lf\t vel_x: %lf\t vel_y: %lf\t
234         brightness: "
235         "%lf\n",
236         p.pos[0], p.pos[1], p.mass, p.vel[0], p.vel[1], p.
237         brightness);
238 }
239
240 /*
241  Function to initialize the root of the the quad tree
242 */
243 q_node_t *init()
244 {
245     q_node_t *root = (q_node_t *)malloc(sizeof(q_node_t));
246
247     //define the initial data
248     root->particle = NULL;
249     root->length = 1.0;
250     root->total_mass = 0.0;
251     root->center_mass = zerovec;
252     root->pos = 0.5 * se_vec;
253
254     //create empty children to the node
255     root->NE = NULL;
256     root->NW = NULL;
257     root->SW = NULL;
258     root->SE = NULL;

```

```

258     return root;
259 }
260
261 /*
262 Recursive insertion of a particle into the quadtree.
263 Updates mass + center of mass of each node, then creates the
264 required child node and inserts
265 particle into it
266 */
267 void insert(q_node_t **node, p_t *particle)
268 {
269     //update mass and center of mass
270     (*node)->center_mass = ((*node)->total_mass * (*node)->
271     center_mass + particle->mass * particle->pos) / ((*node)->
272     total_mass + particle->mass);
273     (*node)->total_mass = (*node)->total_mass + particle->mass;
274
275     //check if a particle is present at the node
276     if ((*node)->particle == NULL)
277     { // if there is no particle, then check if it has children
278
279         if ((*node)->NW == NULL && (*node)->NE == NULL && (*node)->
280         SE == NULL && (*node)->SW == NULL)
281         { // if it doesnt have
282             any children,
283             (*node)->particle = particle; // this means there is
284             no particle or sub-quadrants. Free to add the new particle!!
285             return;
286         }
287         else //insert into the correct child
288         {
289             split_insert(0,particle); //finds the correct child,
290             creates it if not existent, and inserts it into that node
291         }
292     }
293     else
294     {
295         //the particle in the node needs to be inserted into a
296         child instead
297         p_t *old_particle = (*node)->particle;
298         (*node)->particle = NULL;
299
300         // first argument (1) makes if statement in macro function
301         obsolete since if statement
302         // will allways be true
303
304         //insert the particles and create children if needed
305         split_insert(1,particle);
306         split_insert(0,old_particle);
307     }
308 }
309
310 /*
311 Delete tree recursivly, starting at the root
312 */
313 void tree_del(q_node_t **node)

```

```

306 {
307     //check for the existance of children , and delete those that
    are present
308     if ((*node) != NULL) {
309         if ((*node)->NW != NULL)
310             tree_del(&(**node).NW);
311
312         if ((*node)->NE != NULL)
313             tree_del(&(**node).NE);
314
315         if ((*node)->SE != NULL)
316             tree_del(&(**node).SE);
317
318         if ((*node)->SW != NULL)
319             tree_del(&(**node).SW);
320     }
321     free(*node);
322     *node=NULL;
323 }
324
325 /*
326     Traverse down the quad tree and calculate what forces effect a
    particle ,
327     and update its acceleration.
328 */
329 void force_traverse(q_node_t **node, const double theta_max, p_t *
    particle)
330 {
331     //If quadrant is empty or particle is current particle
332     if ((*node) == NULL || (*node)->total_mass == 0 || particle ==
    (*node)->particle)
333     {
334         return;
335     }
336
337     //Get the distance
338     vector_t diff = particle->pos - (*node)->pos;           // particles
    relative position in x,y coordinates
339     vector_t diffsqrd = diff * diff;                         // squared
340     double distance = diffsqrd[0] + diffsqrd[1];             //squared distance
341
342
343     //If theta condition is met calculate acceleration of particle
344     //else traverse down to child nodes
345
346     //NOTE: the original condition is theta_max >= length/distance ,
    but below the
347     //comparison is instead theta_max^2 * distance^2 >= lenght^2.
    Thus we don't need
348     //to take the sqrt for the distance above. ~20% time reduction
    for N=10000 from this
349
350     if (theta_max * distance >= ( (*node)->length) * ((*node)->
    length) ) || (*node)->particle != NULL)
351     {
352         diff = particle->pos - (*node)->center_mass;         //
    particles relative position in x,y coordinates

```

```

353     diffsqrd = diff * diff; //
354     squared
355     distance = sqrt(diffsqrd[0] + diffsqrd[1]) + EPSILON_0; //
356     pythagoras
357     particle->acc -= (*node)->total_mass * diff / (distance *
358     distance * distance);
359     }
360     else
361     {
362         //if theta condition not fulfilled, need to take another
363         //recursive step to each of the children
364         if ((*node)->NW != NULL) force_traverse(&((*node)->NW),
365         theta_max, particle);
366         if ((*node)->NE != NULL) force_traverse(&((*node)->NE),
367         theta_max, particle);
368         if ((*node)->SW != NULL) force_traverse(&((*node)->SW),
369         theta_max, particle);
370         if ((*node)->SE != NULL) force_traverse(&((*node)->SE),
371         theta_max, particle);
372     }
373 }
374
375 /*
376 Barnes Hut algorithm
377 */
378 void barnes_hut(const int nsteps, const double delta_t, p_t *p,
379 const int N, const double theta_max, const int graphics)
380 {
381     // if we use graphics initialize window
382     float circleradius[N], circlecolor = 0;
383     if (graphics) {
384         SetCAxes(0, 1);
385
386         // make radius dependent on the mass of the particle
387         for (register int i = 0; i < N; ++i)
388         {
389             circleradius[i] = fmax(1.5e-3, p[i].mass * 1.5e-3); //have
390             a smallest possible particle size
391         }
392     }
393
394     const double G_dt = (100.0 * delta_t) / N;
395
396     for (int t = 0; t < nsteps; t++)
397     {
398         q_node_t *root = init();
399
400         for (register int i = N - 1; i >= 0; i--)
401         {
402             insert(&root, &(p[i]));
403         }
404
405         for (register int i = N - 1; i >= 0; i--)
406         {
407             force_traverse(&root, theta_max, &(p[i]));
408         }
409     }
410 }

```

```

401     }
402
403     // When all the particles accelerations have been
404     calculated,
405     // we can now, for the same time t_i update all particles
406     velocities
407     // and positions, so that no particle is in a different
408     timestep
409     for (register int i = N - 1; i >= 0; i--)
410     {
411         p[i].vel += G.dt * p[i].acc;
412         p[i].pos += delta_t * p[i].vel;
413         p[i].acc = zerovec;
414     }
415
416     if (graphics) {
417         ClearScreen();
418         // picasso away!
419         for (int i = 0; i < N; i++) {
420             DrawCircle(p[i].pos[0], p[i].pos[1], W, L,
421             circleradius[i], circlecolor);
422         }
423         tree_draw(&root);
424
425         Refresh();
426         usleep(1000);
427     }
428     tree_del(&root);
429 }
430
431 // if we use graphics, flush and close window
432 if (graphics) {
433     FlushDisplay();
434     CloseDisplay();
435 }
436 }
437
438 /*
439     1) recursively traverse the tree, until a region with one
440     particle p_i is reached
441
442     2) traverse into another quadrant
443     - calculate distance r from particle p_i to center of
444     mass of all particles in the quadrant
445     - calculate sidelength h of quadrant (or max distance
446     between any two points in the quadrant)
447     - Theta criterion  $h/r < \theta$ .  $\theta$  usually set to 0.5
448
449     3) CRITERION
450     IF criterion satisfied:
451         - calculate force acting on p_i by the
452         collective shared force of particles in group by center of mass
453     IF criterion NOT satisfied:
454         - traverse recursively into its four quadrants
455         - jump to calculations of 2)
456 */

```

```

450 /*
451     Draws quadrants recursively, starting with the root
452 */
453 void tree_draw(q_node_t **node){
454     if(&(**node) != NULL){
455         DrawRectangle((**node).pos[0]-((**node).length)/2.0, (**
node).pos[1]-((**node).length)/2.0), W, L, (**node).length,
(**node).length, 0.85);
456         tree_draw(&(**node).NW);
457         tree_draw(&(**node).SW);
458         tree_draw(&(**node).NE);
459         tree_draw(&(**node).SE);
460     }
461 }

```

Multithreaded Straightforward galsim.c

```

1  #include <math.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <pthread.h>
5  #include "graphics.h"
6
7  pthread_barrier_t barrier;
8
9  typedef double vector_t __attribute__((vector_size(2 * sizeof(
double))));
10
11 #define EPSILON_0 1e-3
12 int nthreads;
13 const vector_t zerovec = {0.0, 0.0};
14
15 typedef struct particle
16 {
17     /*
18     Particle properties:
19     vel - velocity
20     pos - coordinates
21     */
22     vector_t pos;
23     double mass;
24     vector_t vel;
25     double brightness;
26     vector_t acc;
27 } p_t;
28
29 /*
30 struct that contains all the relevant data for a thread
31 */
32 typedef struct thread_data
33 {
34     int lower_bound;
35     int upper_bound;
36     int thread_id;
37     int N;
38     int nsteps;

```



```

39 double G_dt;
40 double delta_t;
41 p_t *particle_array;
42
43 } thread_data_t;
44
45 void read_doubles_from_file(int N, p_t *p, const char *fileName);
46 void print_struct(p_t p);
47 //void time_sim(int nsteps, const double delta_t, p_t *p, int N,
48 //    const int graphics);
49 void write_to_file(p_t *p, int N, char *fileName);
50 //void time_sim2(int nsteps, const double delta_t, p_t *input_array
51 //    , int N, const int graphics);
52 void time_sim_vec(const int nsteps, const double delta_t, p_t *p,
53     const int N,
54     const int graphics);
55 void *thread_force_update(void *args);
56
57 int main(int argc, char *argv[])
58 {
59     // Check correct number of input arguments
60     if (argc != 7)
61     {
62         printf("Wrong number of inputs \n");
63         exit(EXIT_FAILURE);
64     }
65
66     // Define input variables
67     const int N = atoi(argv[1]);
68     const int nsteps = atoi(argv[3]);
69     const double delta_t = atof(argv[4]);
70     const int graphics = atoi(argv[5]);
71     nthreads = atoi(argv[6]);
72
73     p_t p[N];
74
75     if (graphics)
76     {
77         InitializeGraphics(argv[0], 800, 800);
78         SetCAxes(0, 1);
79         read_doubles_from_file(N, p, argv[2]);
80
81         // print structs if wanted
82         //for (int i=0; i<N; i++) print_struct(p[i]);
83         time_sim_vec(nsteps, delta_t, p, N, graphics);
84         FlushDisplay();
85         CloseDisplay();
86         write_to_file(p, N, "result.gal");
87     }
88     else
89     {
90         read_doubles_from_file(N, p, argv[2]);
91
92         // print structs if wanted
93         //for (int i=0; i<N; i++) print_struct(p[i]);
94         time_sim_vec(nsteps, delta_t, p, N, graphics);
95         write_to_file(p, N, "result.gal");
96     }
97 }

```

```

93     }
94     return 0;
95 }
96
97 void write_to_file(p_t *p, int N, char *fileName)
98 {
99     FILE *file = fopen(fileName, "wb");
100
101     for (int i = 0; i < N; i++)
102     {
103         fwrite(&p[i].pos, sizeof(p[i].pos), 1, file);
104         fwrite(&p[i].mass, sizeof(p[i].mass), 1, file);
105         fwrite(&p[i].vel, sizeof(p[i].vel), 1, file);
106         fwrite(&p[i].brightness, sizeof(p[i].brightness), 1, file);
107     }
108     fclose(file);
109 }
110
111 void print_struct(p_t p)
112 {
113     printf(
114         "x: %lf\t y: %lf\t mass: %lf\t vel_x: %lf\t vel_y: %lf\t\n",
115         p.pos[0], p.pos[1], p.mass, p.vel[0], p.vel[1], p.brightness);
116 }
117
118 ///////////////////////////////////////////////////
119 // READ DOUBLES //
120 ///////////////////////////////////////////////////
121 // NOTE: might not work because vector_t might not operate the same
122 // way as a regular variable.
123
124 void read_doubles_from_file(int N, p_t *p, const char *fileName)
125 {
126     /* Open input file and determine its size. */
127     FILE *file = fopen(fileName, "rb");
128     if (!file)
129     {
130         printf("read_doubles_from_file error: failed to open input file\n",
131             fileName);
132         exit(EXIT_FAILURE);
133     }
134
135     p_t temp;
136     double buffer[6]; // file reading buffer
137     for (int i = 0; i < N; i++)
138     {
139         // Read doubles into particle struct
140         fread(&buffer, sizeof(double), 6, file);
141         temp.pos[0] = buffer[0];
142         temp.pos[1] = buffer[1];
143         temp.mass = buffer[2];
144         temp.vel[0] = buffer[3];
145         temp.vel[1] = buffer[4];

```

```

146     temp.brightness = buffer[5];
147     temp.acc = zerovec;
148     p[i] = temp;
149 }
150 // Close file and check if closed successfully
151 if (fclose(file) != 0)
152 {
153     printf("read_doubles_from_file error: error closing input file
154     .\n");
155     exit(EXIT_FAILURE);
156 }
157
158 // VECTORIZED
159 void time_sim_vec(const int nsteps, const double delta_t, p_t *p,
160                  const int N,
161                  const int graphics)
162 {
163     //vector_t diff, diffsqrd, F;
164
165     //double distance;
166     const double G_dt = delta_t * 100.0 / N;
167
168     float circleradius[N], circlecolor = 0;
169
170     pthread_t threads[nthreads];
171     thread_data_t thread_data[nthreads];
172     pthread_barrier_init(&barrier, NULL, nthreads); //+1 since main
173     // is also waiting
174     pthread_attr_t attr;
175     pthread_attr_init(&attr);
176     pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
177
178     for (int i = 0; i < nthreads; i++)
179     {
180         thread_data[i] = (thread_data_t){.nsteps=nsteps, .lower_bound =
181         i * N / nthreads, .upper_bound = (i + 1) * N / nthreads, .
182         thread_id = i, .G_dt = G_dt, .delta_t = delta_t, .
183         particle_array = p, .N = N};
184         // printf("lbound: %d, ubound %d\n", thread_data[i].lower_bound
185         , thread_data[i].upper_bound);
186     }
187
188     if (graphics)
189     {
190         // make radius dependent on the mass of the particle
191         for (register int i = 0; i < N; ++i)
192         {
193             circleradius[i] = p[i].mass * 1.5e-3;
194         }
195     }
196     for (int thread = 0; thread < nthreads; thread++)
197         pthread_create(&threads[thread], &attr, thread_force_update, &
198         thread_data[thread]);
199 }

```

```

195     for (int thread = 0; thread < nthreads; thread++)
196         pthread_join(threads[thread], NULL);
197
198     pthread_barrier_destroy(&barrier);
199
200     //after all time steps, joint the threads
201 }
202
203 void *thread_force_update(void *args)
204 {
205     thread_data_t *t_data = (thread_data_t *)args;
206     p_t *p = t_data->particle_array;
207     vector_t diff, diffsqrd;
208     double distance;
209
210     for (int t=0; t<t_data->nsteps; t++) {
211         pthread_barrier_wait(&barrier);
212         for (int i = t_data->lower_bound; i < t_data->upper_bound; i++)
213         {
214             for (int j = 0; j < t_data->N; j++) // only up til the particle
215                 in the above foor loop
216             {
217                 if (i == j)
218                     continue;
219                 // temp variables
220                 diff = p[i].pos - p[j].pos; // particles relative position in
221                 x,y coordinates
222                 diffsqrd = diff * diff; // - squared
223                 distance = sqrt(diffsqrd[0] + diffsqrd[1]) + EPSILON_0;
224                 p[i].acc -= p[j].mass * diff / (distance * distance *
225                 distance);
226             }
227         }
228
229         pthread_barrier_wait(&barrier);
230
231         for (int i = t_data->lower_bound; i < t_data->upper_bound; i++)
232         {
233             p[i].vel += t_data->G_dt * p[i].acc;
234             p[i].pos += t_data->delta_t * p[i].vel;
235             p[i].acc = zerovec;
236         }
237     }
238     pthread_exit(NULL);
239 }

```

Multithreaded Barnes-Hut galsim.c

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <pthread.h>
5 #include "graphics.h"
6
7 #define PRINT_TIMES 0 //set to 1 to print execution times of
                        different code blocks

```

```

8
9 #define EPSILON_0 1e-3      //Plummer radius
10 #define W 1                // Window width for graphics
11 #define L 1                // Window length for graphics
12
13
14 /*
15 Globals for Pthread
16 */
17 pthread_barrier_t barrier;
18 int done; //global variable so that a worker thread can know once
           all timesteps are done
19
20
21 typedef double vector_t __attribute__((vector_size(2 * sizeof(
           double))));
22
23 /*
24 Particle:
25 * pos - coordinates
26 * mass - mass of particle
27 * vel - velocity
28 * brightness - brightness of particle
29 * acc - acceleration
30 */
31 typedef struct particle
32 {
33     vector_t pos;
34     double mass;
35     vector_t vel;
36     double brightness;
37     vector_t acc;
38 } p_t;
39
40 /*
41 Quad tree node:
42 * particle - particle struct pointer
43 * length - length of box/quadnode
44 * total_mass - mass of all particles inside a box
45 * center_mass - position of center of mass
46 * pos - coordinate of center of box
47 * NE,NW,SW,SE - quadrants of domain. NE = northeast, etc.. (
           these are the nodes children)
48 */
49 typedef struct quad_tree_node
50 {
51     p_t *particle;
52     double length;
53     double total_mass;
54     vector_t center_mass;
55     vector_t pos;
56     struct quad_tree_node *NE;
57     struct quad_tree_node *NW;
58     struct quad_tree_node *SW;
59     struct quad_tree_node *SE;
60
61 } q_node_t;

```

```

62
63 /*
64 struct that contains all the relevant data for a thread
65 */
66 typedef struct thread_data
67 {
68     uint lower_bound;
69     uint upper_bound;
70     uint thread_id;
71     double theta_max;
72     double G_dt;
73     double delta_t;
74     q_node_t * root;
75     p_t * particle_array;
76 } thread_data_t;
77
78 /*
79     Initialize global helpfull vectors
80 */
81 const vector_t zero_vec = {0.0, 0.0};
82 const vector_t se_vec = {1.0, 1.0};
83 const vector_t sw_vec = {-1.0, 1.0};
84 const vector_t ne_vec = {1.0, -1.0};
85 const vector_t nw_vec = {-1.0, -1.0};
86
87
88
89 /*
90 Macro function that creates a child to the parent, with a vector
91 in the corresponding direction.
92     Example: Create a child to the north-west to a q_node_t ** node
93     create_child((*node)->NW, (*node), nw_vec);
94 */
95 #define create_child(child, parent, dir) \
96 ({ \
97     child = (q_node_t *) malloc(sizeof(q_node_t)); \
98     child->pos = parent->pos + 0.25 * (parent->length) * dir; \
99     child->length = parent->length * 0.5; \
100     child->total_mass = 0.0; \
101     child->particle = NULL; \
102     child->NW = NULL; \
103     child->NE = NULL; \
104     child->SW = NULL; \
105     child->SE = NULL; \
106 })
107
108
109 /*
110 Macro function that locates the correct child quadrant to insert a
111 particle
112 into. If val==0: check if child already exists. If not, create it
113 before inserting.
114 If val==1: insert knowing that child does not exist.
115 */
116 #define split_insert(val, particle) \
117 ({ \
118     if ((particle->pos[0] >= (*node)->pos[0]) && (particle->pos[1]

```

```

117 < (*node)->pos[1])) { \
118     if (val||(*node)->NE == NULL) { \
119         create_child((*node)->NE, (*node), ne_vec); \
120     } \
121     insert(&((*node)->NE), particle); \
122 } \
123 else if ((particle->pos[0] < (*node)->pos[0]) && (particle->pos
124 [1] < (*node)->pos[1])){ \
125     if (val||(*node)->NW == NULL) { \
126         create_child((*node)->NW, (*node), nw_vec); \
127     } \
128     insert(&((*node)->NW), particle); \
129 } \
130 else if ((particle->pos[0] < (*node)->pos[0]) && (particle->pos
131 [1] >= (*node)->pos[1])) { \
132     if (val||(*node)->SW == NULL) { \
133         create_child((*node)->SW, (*node), sw_vec); \
134     } \
135     insert(&((*node)->SW), particle); \
136 } \
137 else { \
138     if (val||(*node)->SE == NULL) { \
139         create_child((*node)->SE, (*node), se_vec); \
140     } \
141     insert(&((*node)->SE), particle); \
142 } \
143 })
144
145 /*
146     Declare functions
147 */
148 //IO functions
149 void read_file(int N, p_t *p, const char *fileName);
150 void write_to_file(p_t *p, int N, char *fileName);
151 void print_struct(p_t p); //mainly for debugging
152 double get_wall_seconds(void); //measuring time
153
154 q_node_t *init(); //create tree root
155 void insert(q_node_t **node, p_t *particle); //recursive insertion
156     of particle into tree
157 void tree_del(q_node_t **node); //free memory recursively
158 void barnes_hut(const int nsteps, const double delta_t, p_t *p,
159     const int N,
160     const double theta_max, const int graphics, const
161     int nthreads); //the time stepping algorithm
162 void *thread_force_update(void *args); //thread force
163 void force_traverse(q_node_t **node, const double theta_max, p_t *
164     particle); //calculation of the forces
165
166 void tree_draw(q_node_t **node); //recursive drawing of the
167     quadnodes

```

```

166  /*
167      Driver function
168  */
169  int main(int argc, char *argv[])
170  {
171      #if PRINT_TIMES
172          double main_time = get_wall_seconds();
173      #endif
174
175
176
177      // Check correct number of input arguments
178      if (argc != 8)
179      {
180          printf("Wrong number of inputs \n");
181          exit(EXIT_FAILURE);
182      }
183
184      // Define input variables
185      const int N = atoi(argv[1]);
186      const int nsteps = atoi(argv[3]);
187      const double delta_t = atof(argv[4]);
188      const double theta_max = atof(argv[5]) * atof(argv[5]); //
189      // square for cheaper comparison in the force calculation
190      const int graphics = atoi(argv[6]);
191      const int nthreads = atoi(argv[7]);
192
193      // Read from input file and save in an array of particles
194      p_t p[N];
195
196      #if PRINT_TIMES
197          double read_time = get_wall_seconds();
198          read_file(N, p, argv[2]);
199          read_time = get_wall_seconds() - read_time;
200          printf("Time spent reading input file: %lf s \n", read_time);
201      #else
202          read_file(N, p, argv[2]);
203      #endif
204
205      // initialize window if graphics is enabled
206      if (graphics) InitializeGraphics(argv[0], 800, 800);
207
208      // Simulate particle system with Barnes Hut algorithm and save
209      // result
210      barnes_hut(nsteps, delta_t, p, N, theta_max, graphics, nthreads);
211
212      #if PRINT_TIMES
213          double write_time = get_wall_seconds();
214          write_to_file(p, N, "result.gal");
215          write_time = get_wall_seconds() - write_time;
216          printf("Time spent saving input file: %lf s \n", write_time);
217      #endif
218
219      main_time = get_wall_seconds() - main_time;

```



```

218     printf("-----\nTotal execution time: %lf s \n", main_time
);
219 #else
220     write_to_file(p, N, "result.gal");
221 #endif
222
223     return 0;
224 }
225
226 /*
227  Function used for measuring times
228 */
229 double get_wall_seconds(void)
230 {
231     struct timeval tv;
232     gettimeofday(&tv, NULL);
233     double seconds = tv.tv_sec + (double)tv.tv_usec / 1000000;
234     return seconds;
235 }
236
237 /*
238  Function that reads doubles from input file and stores them in
  an array of structs
239 */
240 void read_file(int N, p_t *p, const char *fileName)
241 {
242     /* Open input file and determine its size. */
243     FILE *file = fopen(fileName, "rb");
244     if (!file)
245     {
246         printf("read_doubles_from_file error: failed to open input
file '%s'.\n",
247             fileName);
248         exit(EXIT_FAILURE);
249     }
250
251     p_t temp;
252     double buffer[6]; // file reading buffer
253     for (int i = 0; i < N; i++)
254     {
255         // Read doubles into particle struct
256         if (0==fread(&buffer, sizeof(double), 6, file)) exit(
EXIT_FAILURE);
257         temp.pos[0] = buffer[0];
258         temp.pos[1] = buffer[1];
259         temp.mass = buffer[2];
260         temp.vel[0] = buffer[3];
261         temp.vel[1] = buffer[4];
262         temp.brightness = buffer[5];
263         temp.acc = zerovec;
264         p[i] = temp;
265     }
266     // Close file and check if closed successfully
267     if (fclose(file) != 0)
268     {
269         printf("read_doubles_from_file error: error closing input
file.\n");

```

```

270     exit(EXIT_FAILURE);
271 }
272 }
273
274 /*
275  Function that writes output to file
276 */
277 void write_to_file(p_t *p, int N, char *fileName)
278 {
279     FILE *file = fopen(fileName, "wb");
280
281     for (int i = 0; i < N; i++)
282     {
283         fwrite(&p[i].pos, sizeof(p[i].pos), 1, file);
284         fwrite(&p[i].mass, sizeof(p[i].mass), 1, file);
285         fwrite(&p[i].vel, sizeof(p[i].vel), 1, file);
286         fwrite(&p[i].brightness, sizeof(p[i].brightness), 1, file);
287     }
288     fclose(file);
289 }
290
291 /*
292  Function that prints values of a struct. Used for debugging
293 */
294 void print_struct(p_t p)
295 {
296     printf(
297         "x: %lf\t y: %lf\t mass: %lf\t vel_x: %lf\t vel_y: %lf\t\n",
298         p.pos[0], p.pos[1], p.mass, p.vel[0], p.vel[1], p.
299         brightness);
300 }
301
302 /*
303  Function to initialize the root of the the quad tree
304 */
305 q_node_t *init()
306 {
307     q_node_t *root = (q_node_t *)malloc(sizeof(q_node_t));
308
309     //define the initial data
310     root->particle = NULL;
311     root->length = 1.0;
312     root->total_mass = 0.0;
313     root->center_mass = zero_vec;
314     root->pos = 0.5 * se_vec;
315
316     //create empty children to the node
317     root->NE = NULL;
318     root->NW = NULL;
319     root->SW = NULL;
320     root->SE = NULL;
321
322     return root;
323 }
324

```

```

325
326 /*
327 Recursive insertion of a particle into the quadtree.
328 Updates mass + center of mass of each node, then creates the
    required child node and inserts
329 particle into it
330 */
331 void insert(q_node_t **node, p_t *particle)
332 {
333     //update mass and center of mass
334     (*node)->center_mass = ((*node)->total_mass * (*node)->
    center_mass + particle->mass * particle->pos) / ((*node)->
    total_mass + particle->mass);
335     (*node)->total_mass = (*node)->total_mass + particle->mass;
336
337     //check if a particle is present at the node
338     if ((*node)->particle == NULL)
339     { // if there is no particle, then check if it has children
340
341         if ((*node)->NW == NULL && (*node)->NE == NULL && (*node)->
    SE == NULL && (*node)->SW == NULL)
342         { // if it doesnt have
    any children,
343             (*node)->particle = particle; // this means there is
    no particle or sub-quadrants. Free to add the new particle!!
344             return;
345         }
346         else //insert into the correct child
347         {
348             split_insert(0,particle); //finds the correct child,
    creates it if not existent, and inserts it into that node
349         }
350     }
351     else
352     {
353         //the particle in the node needs to be inserted into a
    child instead
354         p_t *old_particle = (*node)->particle;
355         (*node)->particle = NULL;
356
357         // first argument (1) makes if statement in macro function
    obsolete since if statement
358         // will allways be true
359
360         //insert the particles and create children if needed
361         split_insert(1,particle);
362         split_insert(0,old_particle);
363     }
364 }
365
366 /*
367 Delete tree recursivly, starting at the root
368 */
369 void tree_del(q_node_t **node)
370 {
371     //check for the existance of children, and delete those that
    are present

```

```

372     if ((*node) != NULL) {
373         if ((*node)->NW != NULL)
374             tree_del(&(**node).NW);
375
376         if ((*node)->NE != NULL)
377             tree_del(&(**node).NE);
378
379         if ((*node)->SE != NULL)
380             tree_del(&(**node).SE);
381
382         if ((*node)->SW != NULL)
383             tree_del(&(**node).SW);
384     }
385     free(*node);
386     *node=NULL;
387 }
388
389 /*
390  Traverse down the quad tree and calculate what forces effect a
391  particle ,
392  and update its acceleration .
393 */
394 void force_traverse(q_node_t **node, const double theta_max, p_t *
395 particle)
396 {
397     //If quadrant is empty or particle is current particle
398     if ((*node) == NULL || (*node)->total_mass == 0 || particle ==
399 (*node)->particle)
400     {
401         return;
402     }
403
404     //Get the distance
405     vector_t diff = particle->pos - (*node)->pos; // particles
406             relative position in x,y coordinates
407     vector_t diffsqrd = diff * diff; // squared
408     double distance = diffsqrd[0] + diffsqrd[1]; //squared distance
409
410     //If theta condition is met calculate acceleration of particle
411     //else traverse down to child nodes
412
413     //NOTE: the original condition is theta_max >= length/distance ,
414     but below the
415     //comparison is instead theta_max^2 * distance^2 >= lenght^2.
416     Thus we don't need
417     //to take the sqrt for the distance above. ~20% time reduction
418     for N=10000 from this
419
420     if (theta_max * distance >= ( (*node)->length) * ((*node)->
421 length) ) || (*node)->particle != NULL)
422     {
423         diff = particle->pos - (*node)->center_mass; //
424         particles relative position in x,y coordinates
425         diffsqrd = diff * diff; //
426         squared
427         distance = sqrt(diffsqrd[0] + diffsqrd[1]) + EPSILON_0; //

```

```

419     pythagoras
420         particle->acc -= (*node)->total_mass * diff / (distance *
421         distance * distance);
422     }
423     else
424     {
425         //if theta condition not fulfilled, need to take another
426         //recursive step to each of the children
427         if ((*node)->NW != NULL) force_traverse(&((*node)->NW),
428         theta_max, particle);
429         if ((*node)->NE != NULL) force_traverse(&((*node)->NE),
430         theta_max, particle);
431         if ((*node)->SW != NULL) force_traverse(&((*node)->SW),
432         theta_max, particle);
433         if ((*node)->SE != NULL) force_traverse(&((*node)->SE),
434         theta_max, particle);
435     }
436 }
437
438 /*
439 Barnes Hut algorithm
440 */
441 void barnes_hut(const int nsteps, const double delta_t, p_t *p,
442 const int N, const double theta_max, const int graphics, const
443 int nthreads)
444 {
445     #if PRINT_TIMES //declare variables if we want ot emasure times
446     double build_time=0.0;
447     double force_time=0.0;
448     double delete_time=0.0;
449     #endif
450
451     // if we use graphics initialize window
452     float circleradius[N], circlecolor = 0;
453     if (graphics) {
454         SetCAxes(0, 1);
455
456         // make radius dependent on the mass of the particle
457         for (register int i = 0; i < N; ++i)
458         {
459             circleradius[i] = fmax(1.5e-3, p[i].mass * 1.5e-3); //have
460             a smallest possible particle size
461         }
462     }
463
464     const double G_dt = (100.0 * delta_t) / N;
465
466     pthread_t threads[nthreads];
467     thread_data_t thread_data[nthreads];
468     pthread_barrier_init(&barrier, NULL, nthreads+1); //+1 since
469     main is also waiting
470
471

```

```

466 q_node_t *root = init();
467 done = 0; //0 means we are not done with simulation
468
469
470 for (int i=0; i<nthreads; i++) {
471     thread_data[i] = (thread_data_t) { .lower_bound = i*N/
472     nthreads, .upper_bound=(i+1)*N/nthreads,
473     .thread_id=i, .theta_max=theta_max, .G_dt=
474     G_dt, .delta_t=delta_t,
475     .root = root, .particle_array=p};
476     // printf("lbound: %d, ubound %d\n", thread_data[i].
477     lower_bound, thread_data[i].upper_bound);
478 }
479
480 for (int thread=0; thread<nthreads; thread++) {
481     pthread_create(&threads[thread], NULL,
482     thread_force_update, &thread_data[thread]);
483 }
484
485 for (register int t = 0; t < nsteps; t++)
486 {
487     #if PRINT_TIMES
488     build_time -= get_wall_seconds();
489     for (register int i = N - 1; i >= 0; i--) insert(&root,
490     &(p[i]));
491     build_time += get_wall_seconds();
492     #else
493     for (register int i = N - 1; i >= 0; i--) insert(&root,
494     &(p[i]));
495     #endif
496
497     pthread_barrier_wait(&barrier); //unlocks barrier to start
498     time steps in threads
499
500     #if PRINT_TIMES
501     force_time -= get_wall_seconds();
502     pthread_barrier_wait(&barrier);
503     force_time += get_wall_seconds();
504     #else
505     pthread_barrier_wait(&barrier);
506     #endif
507
508     if (graphics) {
509         ClearScreen();
510         // picasso away!
511         for (int i = 0; i < N; i++) {
512             DrawCircle(p[i].pos[0], p[i].pos[1], W, L,
513             circleradius[i], circlecolor);
514         }
515         tree_draw(&root);
516
517         Refresh();
518         usleep(3000);

```

```

515     }
516
517     // tree_del(&root);
518
519     //TEMP BLOCK
520     #if PRINT_TIMES
521         delete_time -= get_wall_seconds();
522         tree_del(&(root->NE)); tree_del(&(root->NW)); tree_del
523         (&(root->SW)); tree_del(&(root->SE));
524         delete_time += get_wall_seconds();
525     #else
526         tree_del(&(root->NE)); tree_del(&(root->NW)); tree_del
527         (&(root->SW)); tree_del(&(root->SE));
528     #endif
529
530     //RESET ROOT: SAME AS INIT BUT NO MALLOC. THIS WAY WE CAN
531     KEEP THE ADDRESS TO ROOT
532     //AND NOT CHANGE IT EACH TIME STEP
533     root->particle = NULL;
534     root->total_mass = 0.0;
535     root->center_mass = zero_vec;
536     root->NE = NULL;
537     root->NW = NULL;
538     root->SW = NULL;
539     root->SE = NULL;
540
541     }
542
543     done=1;
544     pthread_barrier_wait(&barrier);
545     free(root);
546
547     // //after all time steps, joint the threads
548     for (int thread=0; thread<nthreads; thread++) pthread_join(
549     threads[thread], NULL);
550     pthread_barrier_destroy(&barrier); //destroy barrier once it is
551     no longer used
552     // if we use graphics, flush and close window
553     if (graphics) {
554         FlushDisplay();
555         CloseDisplay();
556     }
557
558     #if PRINT_TIMES
559         printf("Time spent building the quadtree: %lf s\n",
560         build_time);
561         printf("Time spent calculating forces and updating
562         particles: %lf s\n", force_time);
563         printf("Time spent on deleting the quadtree: %lf s\n",
564         delete_time);
565     #endif
566 }
567
568 void *thread_force_update(void * args) {

```

```

564     thread_data_t * t_data = (thread_data_t *) args;
565     p_t * p = t_data->particle_array;
566
567     while( True ) {
568         pthread_barrier_wait(&barrier); //wait until all threads
569         are in position to start time step (tree built)
570         if (done) break; //we set done to 1 in barnes-hut() after
571         all time steps complete
572
573         for (uint i = t_data->lower_bound; i<t_data->upper_bound; i
574         ++){
575             force_traverse(&(t_data->root), t_data->theta_max,
576             &(p[i]));
577             p[i].vel += t_data->G_dt * p[i].acc;
578             p[i].pos += t_data->delta_t * p[i].vel;
579             p[i].acc = zerovec;
580         }
581
582         pthread_barrier_wait(&barrier); //wait until all
583         threads done (s.t. master can start deleting)
584     }
585
586 }
587
588 /*
589  Draws quadrants recursively, starting with the root
590 */
591 void tree_draw(q_node_t **node){
592     if(&(**node) != NULL){
593         DrawRectangle(((**node).pos[0]-(((**node).length)/2.0), (**
594         node).pos[1]-(((**node).length)/2.0), W, L, (**node).length,
595         (**node).length, 0.85);
596         tree_draw(&(**node).NW);
597         tree_draw(&(**node).SW);
598         tree_draw(&(**node).NE);
599         tree_draw(&(**node).SE);
600     }
601 }

```